

Apollo

A governed engine for designing with AI — on brand, accessible, at speed.

What you will see in the next sixteen minutes

#	MINS	CLOCK	PLOT POINT	ON SCREEN
B1	2:00	0:00–2:00	INTRO	
01	1:00	0:00–1:00	Why I am here, in one sentence	Title card
02	1:00	1:00–2:00	What you will see in the next sixteen minutes	The strip above
B2	4:00	2:00–6:00	GOAL · PRACTICAL BENEFITS · PREDICTED KPIS	
03	1:30	2:00–3:30	The problem, in one metaphor	Title card — metaphor spoken, no diagram
04	1:00	3:30–4:30	The claim: two measured gains, stacked — and the hypothesis the pilot tests	Number card — 15% · 47% · 6 mo → 6 wk?
05	1:30	4:30–6:00	The only measured before-and-after we own	The cold-run scorecard
B3	5:30	6:00–11:30	RUN THE DEMO — THE FROZEN PROMPT, COLD	
06	0:30	6:00–6:30	The prompt goes in, unedited	The editor, the released skill pack
07	2:00	6:30–8:30	Dead air A — what inspired it	The build scrolling. Do not switch away.
08	1:30	8:30–10:00	Dead air B — how it developed	Timeline card, if the build is still running
09	1:30	10:00–11:30	The page lands. Drive it.	The generated dashboard, light theme
B4	4:30	11:30–16:00	THE SYSTEM, SIMPLY	
10	1:30	11:30–13:00	Four moving parts, no black box	Four-part diagram
11	3:00	13:00–16:00	The knowledge graph — dig, scrub, trace, and ask it why	The knowledge-graph explorer
B5	2:00	16:00–18:00	AMBITION · REQUIREMENTS · THE ASK	
12	2:00	16:00–18:00	Where it goes, what it needs, what I am asking for	The ask card, then silence
	18:00		Stated sum. Band 15–20. Two minutes of deliberate slack.	

— PLOT POINT 03 · THE PROBLEM, IN ONE METAPHOR

The parts bin

An assembly line is only as fast as its parts bin — and our designers have been hand-machining the gearbox every time a car comes down the line.

Spoken, not drawn. Source: the run of show.

Two measured gains, stacked — and **the hypothesis the pilot tests.**

15%

A major bank's own published figure

Coding-time efficiency across
20,000+ developers.

47%

IBM Carbon, no AI

Same form, 4.2 h → 2.0 h with a
design system. The library alone.

6 months
→ 6 weeks?

Hypothesis, measured by the pilot

Not a result. "How do you know six
weeks?" — "I don't. That is K1."

Sources: a major bank's published AI-efficiency figure · Sparkbox's IBM Carbon study.

— PLOT POINT 05 · THE ONLY MEASURED BEFORE-AND-AFTER WE OWN

BEFORE — 8 SEPTEMBER, FIRST RUN



Four axes — visually rich, full, persistent, interactive — nought to three each. Five blind runs in one day, each one finding a real defect. Self-scored, and every point verified by driving the page in a browser, not by reading the source.

Source: our own cold-run scorecard, 8 September 2026.

— PLOT POINT 05 · PREDICTED KPIS — FIVE, EACH WITH AN EXTERNAL BASELINE

	NAME	BASELINE	TARGET ON THE FIRST LIVE PROJECT
K1	Time to first gated screen	Sparkbox 4.2 h → 2.0 h	Same-day for a multi-screen flow — the only KPI the demo produces live
K2	Library reuse rate	GitClear: AI duplication up 8×	≥80% of rendered UI from gated components
K3	Change-failure rate	DORA 2025: AI adoption correlates with worse stability	No worse than the team's pre-AI rate — explicitly not an improvement
K4	Research-analysis cycle time	JMIR 2024: 492 → 15 min, 71–90% consistency	≥70% on the analysis step; fieldwork hours unchanged and reported separately
K5	Coding-time efficiency	The bank's own published 15%	Match it on the front-end build. Promise no more.

Sources: published industry studies. Designer-hours, headcount and “×” multipliers are deliberately excluded.

— BEAT 3 · RUN THE DEMO

Live build.

The frozen prompt, cold. One prompt, one compliant dashboard,
nothing edited.

— SWITCH TO THE EDITOR

Live and unedited: the frozen prompt, run cold.

Four moving parts. No black box.

01 Canon The design tokens and the reviewed, gated components — the single source of truth for what “on brand” means.	02 Gates Executable checks for brand fidelity and accessibility. Not guidance — enforcement that can withhold “done”.	03 Runbooks The method: how a task is generated, retrieved, and reviewed. Repeatable, inspectable, improvable.	04 Host agent The AI that operates it — running today in the editor designers already have.
--	--	---	--

Portable by design — no bespoke runtime. Source: the Apollo sponsor briefing, July 2026.

— PLOT POINT 11 · THE KNOWLEDGE GRAPH — DIG, SCRUB, TRACE, AND ASK IT WHY

The graph is the source of truth.

Live graph — 2,837 nodes · 4,902 edges. The same field as the title card, only now you can drive it.

—— SWITCH TO THE KNOWLEDGE-GRAPH EXPLORER

Source: the live Apollo knowledge graph, 11 September 2026.

One live project, sponsored — six weeks on the clock, K1–K5 measured against your own published fifteen percent. If it doesn't beat your number, you've lost six weeks. If it does, you own the stack.

Draft wording. The ask is not settled — amount and shape are still open.

— IN CLOSING

Apollo

The build is automated.
The craft is the point.

2,837 nodes · 4,902 edges